

NON-INTRUSIVE LOAD MONITORING FOR HOUSEHOLD APPLIANCE ENERGY DISAGGREGATION

Student Name	University ID
Saif Eddin Jamil Mohammad Mhaisen	0223489
Hamza Ziad Mustafa Ali	2213243
Bashar Nazeh Youssef Zabaneh	2221283
Obada Mo'tasem Yaseen Sultan	2221530

Project Supervisor: Dr. Ruba Obiedat



Department of Computer Science
KING ABDULLAH II SCHOOL OF INFORMATION TECHNOLOGY
THE UNIVERSITY OF JORDAN
May 2026

ABSTRACT

Non-Intrusive Load Monitoring (NILM) is the task of disaggregating household appliance-level energy consumption from a single whole-home smart meter signal. Rather than installing individual sensors on every device, NILM uses machine learning to infer which appliances are active and how much power each consumes — using only the aggregate reading from a single meter at the home's entry point.

This project implements a complete end-to-end NILM system covering four appliances from the REFIT dataset: Kettle, Fridge, Washing Machine, and Dishwasher. Each appliance has a dedicated deep learning model based on the Seq2Point windowing convention, reformulated as binary classification. Development proceeded through three phases: (1) a classical machine learning baseline using SVM and Random Forest with 16 hand-crafted features, achieving F1 up to 0.548; (2) an initial BiLSTM phase achieving F1=0.67 on the kettle; and (3) a final CNN-BiLSTM phase with per-appliance window size and threshold tuning.

All models were evaluated under strict Leave-One-House-Out (LOHO) protocol using House 20 as the completely unseen validation house. Final F1 scores: Kettle 0.784, Fridge 0.688, Washing Machine 0.577, Dishwasher 0.726. The system is deployed as a five-page Streamlit web application providing disaggregation, analytics, bill estimation using Jordan EMRC tariffs, energy forecasting, and a savings simulator.

Keywords: *Non-Intrusive Load Monitoring, NILM, Energy Disaggregation, Deep Learning, CNN-BiLSTM, Seq2Point, REFIT Dataset, Smart Meter.*

ACKNOWLEDGEMENTS

We would like to express our sincere gratitude to our supervisor, Dr. Ruba Obiedat, for her continuous guidance, invaluable feedback, and encouragement throughout both Project-1 and Project-2. Her direction on the evaluation protocol and model architecture was instrumental in shaping the final system.

We are grateful to the King Abdullah II School of Information Technology at the University of Jordan for providing the academic environment and resources that made this project possible.

We also acknowledge the creators of the REFIT Electrical Load Measurements dataset, whose open data enabled the research conducted in this project.

LIST OF ABBREVIATIONS

Abbreviation	Definition
BCE	Binary Cross-Entropy Loss
BiLSTM	Bidirectional Long Short-Term Memory
CNN	Convolutional Neural Network
EMRC	Electricity and Minerals Regulatory Commission (Jordan)
F1	F1 Score — harmonic mean of Precision and Recall on the ON class
GPU	Graphics Processing Unit
KASIT	King Abdullah II School of Information Technology
LOHO	Leave-One-House-Out evaluation protocol
LSTM	Long Short-Term Memory
MAE	Mean Absolute Error
NILM	Non-Intrusive Load Monitoring
REFIT	Residential Energy Disaggregation Dataset (Murray et al., 2017)
RF	Random Forest
SVM	Support Vector Machine
UJ	The University of Jordan
WM	Washing Machine

CHAPTER ONE: INTRODUCTION

1.1 Project Motivation

Non-Intrusive Load Monitoring (NILM) is the task of disaggregating household appliance-level energy consumption from a single whole-home smart meter signal. Rather than installing individual sensors on every device, NILM uses machine learning to infer which appliances are active and how much power each consumes — using only the aggregate reading from a single meter at the home's entry point.

This project implements a complete end-to-end NILM system covering four appliances from the REFIT Electrical Load Measurements dataset: Kettle, Fridge, Washing Machine, and Dishwasher. Each appliance has a dedicated deep learning model, evaluated using strict leave-one-house-out (LOHO) protocol with House 20 as the unseen test house, and deployed through a fully functional Streamlit web application.

1.2 System Architecture

Component	Details
Dataset	REFIT Electrical Load Measurements — 20 UK households at 8-second intervals
Resolution	Resampled to 1-minute intervals (~7× reduction from raw 8-sec data)
Training Houses	2, 3, 5, 9, 11 (varies by appliance availability)
Validation House	House 20 — completely unseen during all training and tuning
Evaluation Protocol	Leave-One-House-Out (LOHO) — strictest standard in NILM research
Primary Metric	F1 Score on ON class, evaluated on full 623,419-window validation set
Framework	PyTorch 2.6 on Google Colab + Google Drive
Deployment	Streamlit web application (nilm_v2)

1.3 Target Appliances — Final Results Summary

Appliance	Final Model	Architecture	F1	Precision	Recall	Threshold
Kettle	kettle_v8_w99.pt	Pure CNN, w=99	0.784	0.726	0.853	0.85
Fridge	fridge_v4_best.pt	CNN- BiLSTM, w=99	0.688	0.545	0.933	0.22
Washing Machine	washing_machine_v5_final.pt	CNN- BiLSTM, w=599	0.577	0.561	0.593	0.80
Dishwasher	dishwasher_multiscale_best.pt	4-ch MultiScale CNN- BiLSTM, w=599	0.726	0.706	0.747	0.80

Note on Microwave: The microwave was trained extensively (best result: F1=0.40 using DilatedCNN + Focal Loss at 8-second resolution). This was accepted as the practical ceiling and microwave was excluded from the final four-appliance lineup.

1.4 How Our Approach Compares to Prior Work

Aspect	Liao et al. 2014 (IEEE)	Kelly & Knottenbelt 2015 (UK-DALE)	Our Project
Dataset	Various household datasets	UK-DALE dataset	REFIT dataset (20 UK households)
Evaluation	Within-house splits	Mixed train/test splits	Leave-one-house-out (strictest)
Architecture	Classical ML / HMM	LSTM Seq2Seq (regression)	CNN-BiLSTM Seq2Point (classification)
Task Framing	Regression (watt estimation)	Regression (watt estimation)	Binary classification + mean-watts estimate
Window	Fixed segments	Sliding window 599	Per-appliance tuned (99 or 599 min)
Validation	Seen houses	Partially seen houses	Completely unseen House 20
MultiScale Input	No	No	Yes (dishwasher: 4-channel input)

The key advancement of our work is the combination of (1) per-appliance window size tuning — which we show is more impactful than architecture changes — and (2) honest evaluation on a completely unseen house under strict LOHO protocol. Our results are not inflated by evaluating on houses seen during training.

1.5 Software and Hardware Requirements

Training was conducted on Google Colab using T4 GPU acceleration with PyTorch 2.6. Data was stored on Google Drive at /content/drive/MyDrive/new obada/. The Streamlit application requires Python 3.10+, PyTorch 2.6, Streamlit 1.30+, pandas, and NumPy. No additional hardware beyond the smart meter is required for deployment.

1.6 Report Outline

Chapter 2 documents the classical ML baseline phase (SVM, Random Forest, initial BiLSTM). Chapter 3 describes the REFIT dataset and the preprocessing pipeline. Chapter 4 presents all Kettle experiments and results. Chapter 5 covers the Fridge. Chapter 6 covers the Washing Machine. Chapter 7 covers the Dishwasher. Chapter 8 synthesises cross-cutting lessons. Chapter 9 describes the Streamlit application and real-world testing. Chapter 10 compares final results to published literature. Chapter 11 presents conclusions.

CHAPTER TWO: CLASSICAL ML & EARLY EXPERIMENTS

Before arriving at the final deep learning architecture, the project passed through a complete classical machine learning phase. These early experiments — SVM, Random Forest, and an initial Bidirectional LSTM — were essential for establishing baseline performance, understanding the difficulty of the problem, and justifying the pivot to a dedicated per-appliance CNN-BiLSTM Seq2Point approach.

2.1 Setup & Data

All classical ML experiments used the same REFIT dataset, with Houses 2, 3, 5, 9, and 11 for training, and House 20 as the validation house. Code ran on Google Colab (T4 GPU). Each model received 16 hand-crafted features extracted from a sliding window over the aggregate signal: mean, std, min, max, peak-to-peak, energy, zero-crossing rate, and several differential statistics.

Parameter	Value
Training Houses	Houses 2, 3, 5, 9, 11
Validation House	House 20 (completely unseen)
Features per Window	16 (mean, std, min, max, peak-to-peak, energy, zero-crossing rate, differentials)
Compute Environment	Google Colab T4 GPU + Google Drive
Key Scripts	svm_colab.py, rf_colab.py, lstm_final.py

2.2 SVM (Support Vector Machine)

An SVM with RBF kernel was trained separately per appliance. Features were standardised using StandardScaler. Class weighting was set to 'balanced' to handle ON/OFF imbalance.

Appliance	Val F1	Val Precision	Val Recall
Kettle	0.000	—	—
Fridge	0.364	—	—

Appliance	Val F1	Val Precision	Val Recall
Washing Machine	0.000	—	—
Dishwasher	0.002	—	—
Microwave	0.000	—	—

Why SVM Failed: The SVM is a margin-based classifier that cannot capture temporal patterns — it sees 16 aggregate statistics but has no concept of the time sequence of power. Appliances like the kettle operate for only 0.3% of all minutes; the SVM collapsed to always predicting OFF. The single-power-threshold limitation meant the model could not distinguish overlapping signatures when multiple appliances ran simultaneously.

2.3 Random Forest

A Random Forest classifier (300 estimators, max_depth=20, min_samples_leaf=5, class_weight='balanced') was trained using the same 16 features but without StandardScaler. This was the best-performing classical ML approach overall.

Appliance	Val F1	Val Prec	Val Rec	Test F1	Test Prec	Test Rec
Kettle	0.045	0.023	0.783	0.057	0.029	0.820
Fridge	0.497	0.417	0.615	0.408	0.521	0.335
Washing Machine	0.548	0.438	0.729	0.499	0.382	0.723
Dishwasher	0.468	0.318	0.884	N/A	N/A	N/A
Microwave	0.050	0.032	0.116	0.041	0.029	0.069

Why RF Was Better Than SVM: Random Forest handles non-linear decision boundaries better and does not require feature scaling. Despite this, the fundamental limitation remained: all 16 features describe the aggregate signal statistically but discard the temporal sequence. High recall (kettle 0.783, washing machine 0.729) was paired with very low precision, indicating many false positives.

2.4 Initial LSTM / BiLSTM (Seq2Point — First Attempt)

The first deep learning attempt was a Seq2Point regression model. It was abandoned due to severe training instability (prediction ratio oscillating $0\times-15\times$ between epochs) and replaced by a BiLSTM classification approach.

Experiment	Result / Notes
Seq2Point multitask (regression)	ABANDONED — oscillating prediction ratio; supervisor indicated Seq2Point was outdated
BiLSTM classification (Kettle)	Val F1=0.67, Test F1=0.63 — first competitive result; foundation for deep learning phase
Key bug: RAM crash	Fixed by chunked streaming extraction (50K rows/chunk)
Key bug: zero collapse	Fixed by balanced 50/50 ON/OFF batch construction
Key bug: IndexError	Fixed by deriving labels from raw power files at runtime

2.5 Phase 0 Model Comparison Summary

Appliance	SVM Best F1	RF Best F1	LSTM Best F1	Phase 0 Winner
Kettle	0.000	0.045	0.67	BiLSTM
Fridge	0.364	0.497	—	Random Forest
Washing Machine	0.000	0.548	—	Random Forest
Dishwasher	0.002	0.468	—	Random Forest
Microwave	0.000	0.050	—	None (all fail)

Root cause of classical ML failure: Both SVM and Random Forest receive a fixed vector of 16 aggregate statistics and must classify from that alone. They cannot learn temporal patterns — the rhythm of a fridge compressor cycling every 60 minutes, the multi-phase signature of a washing machine, or the brief 1-minute spike of a kettle. The BiLSTM's ability to process the full sequence in both directions was the fundamental breakthrough that classical methods could not replicate.

CHAPTER THREE: DATASET & PREPROCESSING PIPELINE

3.1 REFIT Dataset Overview

Property	Value
Houses available	20 (Houses 1–21; House 14 missing)
Sampling interval	~8 seconds (raw)
Project resolution	1 minute (mean-resampled)
Appliance columns per house	Up to 9 (Appliance1–Appliance9; mapping varies by house)
Column mapping source	MetaData_Tables.me.xlsx (one sheet per house)
Meter map output	confirmed_meters.json (all houses, all target appliances)

3.2 House Selection

House	Fridge	Kettle	Microwave	Washing Machine	Dishwasher	Role
2	App1	App8	App5	App2	App3	Training
3	App2	App9	App8	App6	App5	Training
5	App1	App8	App7	App3	App4	Training
9	App1	App7	App6	App3	App4	Training
11	App1	App7	App6	App3	App4	Training
20	App1	App9	App8	App4	App5	VALIDATION ONLY

Houses excluded: House 4 (two washing machines — ambiguous mapping), House 8 (Washer Dryer vs. Washing Machine conflict), House 1 (no Kettle or Microwave), House 6 (no Fridge).

3.3 Appliance Configurations

Appliance	ON Threshold	Min ON Duration	Training Stride	ON% in All Data
Kettle	2000 W	1 min	1	0.5%
Microwave	200 W	1 min	1	0.7%
Fridge	50 W → 30 W*	2 min	5	41.2%
Washing Machine	20 W	5 min	2	4.4%
Dishwasher	10 W	5 min	2	5.4%

* Fridge ON threshold lowered from 50 W to 30 W after diagnostic — House 20 fridge median was 0 W even when labelled ON at 50 W.

3.4 Pipeline Stages

Cell	Purpose	Key Output
1	Mount Google Drive	Drive connected
2	Configuration	All thresholds, paths, house lists defined centrally
3	Helper Functions	<code>load_and_resample()</code> , <code>fill_gaps()</code> , <code>compute_labels()</code> , <code>extract_windows_array()</code>
4	Labelled CSVs	One CSV per appliance (all houses combined) with checkpoint system
5	PKL Windows	Chunked sliding window extraction (50K rows/chunk) — RAM-safe
6	Sanity Check	Verify shapes, ON/OFF ratios, file sizes for all outputs

3.5 Key Design Decisions

Resampling to 1 minute reduces data $7\times$ and aligns with NILM literature. Gap handling: forward-fill ≤ 10 minutes; larger gaps excluded at window extraction. Minimum ON duration filter removes spurious sensor spikes. Training batches are always 50/50 balanced ON/OFF; validation is always left at the natural unbalanced distribution. A checkpoint system ensures every cell skips already-completed work, making the pipeline safe against Colab disconnections.

3.6 Phase 1 Output Summary

Appliance	Total Rows	ON Minutes	Train Windows	Val Windows	Val ON%
Kettle	4,139,248	22,423 (0.5%)	40,554	623,419	0.3%
Microwave	4,139,248	27,944 (0.7%)	49,428	623,419	0.5%
Fridge	4,139,248	1,707,333 (41.2%)	579,342	623,419	41.3%
Washing Machine	4,139,248	180,916 (4.4%)	165,692	623,419	2.4%
Dishwasher	4,139,248	224,663 (5.4%)	215,858	623,419	1.4%

CHAPTER FOUR: KETTLE — EXPERIMENTS & RESULTS

4.1 Challenge

Despite having the strongest power signature (2,000–3,000 W), the kettle is difficult due to extreme temporal sparsity: only 0.3% of all minutes are ON. A model predicting always-OFF achieves 99.7% accuracy — making accuracy a completely useless metric. Every decision had to be evaluated against F1 on the real unbalanced distribution.

4.2 All Experiments

Version	Architecture	Window	Loss	F1	Notes
v1	CNN-BiLSTM	599	BCE	0.63	Double sigmoid bug — loss stuck at 0.52
v2	CNN-BiLSTM	599	Focal	0.60	Fixed sigmoid; slight improvement
v3	CNN-BiLSTM	599	Focal	0.67	Best at window=599; 1:10 imbalance
v4	CNN-BiLSTM	599	Focal	0.63	Lower LR — worse than v3
v5	CNN-BiLSTM	599	Focal	0.43	Event-centered windowing — failed
v6	CNN-BiLSTM	599	Focal	0.44	7 houses added — failed
v7	CNN-BiLSTM	599	MSE	0.52	Regression — bimodal signal unsuitable
v8 ★	Pure CNN	99	Focal	0.784	WINNER — window change + CNN alone
v9	Pure CNN	49	Focal	0.79	Similar to v8, higher precision/lower recall

4.3 Key Failures

Double Sigmoid Bug (v1): `nn.Sigmoid()` inside model + `BCEWithLogitsLoss` both apply sigmoid. Probabilities collapse to 0.5–0.73. Fix: remove sigmoid from model entirely.

Event-centered windowing (v5): ± 8 minute offsets placed window centres outside the 1-minute ON event. Model never saw any training event. F1=0.43.

Regression (v7): kettle signal is bimodal (0W or 2500W). `MSELoss` converges to the mean (~ 12 W) — wrong for both classes.

4.4 Final Model — kettle_v8_w99.pt

Parameter	Value
Architecture	4-layer Pure CNN (no BiLSTM)
Window size	99 minutes
Training imbalance	1:10 (ON:OFF)
Loss	FocalLoss(alpha=0.75, gamma=2.0)
Parameters	95,617
F1 / Precision / Recall	0.784 / 0.726 / 0.853
Threshold	0.85
Power estimate	2,478 W mean ON watts

CHAPTER FIVE: FRIDGE — EXPERIMENTS & RESULTS

5.1 Challenge

The fridge compressor draws only 50–100 W against a household aggregate mean of 376 W — giving a signal-to-noise ratio of approximately 13%. The model must detect a 50 W change in a 376 W noisy background, where the change occurs every 10–20 minutes and lasts 20–40 minutes.

5.2 All Experiments

Version	Architecture	Window	Full Val F1	Separation	Outcome
V1	CNN-BiLSTM (3-layer)	599	0.507	0.078	Separation too low
V2	Dilated CNN + Focal + Attention	599	0.528	0.077	Architecture NOT the bottleneck
V3a	CNN-BiLSTM + rolling std + diff	599	0.484	0.067	Feature engineering made it WORSE
V3b	2D CNN on Spectrogram	599	0.67 (fake)	—	Collapsed: Recall=1.0, predicts ON always
V4 ★	CNN-BiLSTM, smaller	99	0.688	0.296	WINNER — 4× improvement in separation
V5	Deeper CNN-BiLSTM + Focal alpha=0.4	99	0.68	—	Overfitting from epoch 6
V6	CNN-BiLSTM + hour + day features	99	0.690	0.287	Negligible gain — not worth 3× storage

Window Size Breakthrough: Reducing the window from 599 minutes (10 hours) to 99 minutes (1.5 hours) was the single most impactful change. At window=599, the model sees 30–60 complete cycles, making any single cycle invisible. At window=99, the compressor rhythm (5–10 cycles) becomes a learnable pattern. ON/OFF separation improved from 0.078 to 0.296 — a 4× improvement — with no architecture change.

5.3 Final Model — fridge_v4_best.pt

Parameter	Value
Architecture	CNN-BiLSTM (3 CNN layers + 2-layer BiLSTM)
Window size	99 minutes
Training houses	2, 3, 5, 9 (House 11 removed — broken sensor, ON ratio = 0.002)
Parameters	199,489
F1 / Precision / Recall	0.688 / 0.545 / 0.933
Threshold	0.22
ON/OFF separation	0.296 (vs 0.078 at window=599)
Power estimate	85 W mean ON watts

CHAPTER SIX: WASHING MACHINE — EXPERIMENTS & RESULTS

6.1 Challenge

Three compounding factors make the washing machine the hardest appliance for energy estimation: 2.4% ON ratio in House 20 (severe imbalance), extremely high power variance across cycle phases (std=699W, range 20W–2,400W), and house-specific usage patterns.

6.2 Classification Experiments

Experiment	Key Change	Full Val F1	Outcome
v1 — BCE + pos_weight=40	Weighted loss for 40:1 imbalance	0.667 (all ON)	FAILED — zero collapse from epoch 1
v2 — pos_weight=5	Reduced weight	Unstable	Better but inflated val sample metric
v3 — Focal + 1:6 val	FocalLoss + realistic val ratio	0.559	More stable but still low
v4 — Houses 5,9,11 only	Removed weak houses 2 & 3	0.581	Small gain but lost 74K training windows
v5 — Smart Sampling ★	Cap weak houses at 5K ON; keep strong full	0.600	Best classification result
MultiScale 4-ch v2	4-channel input + OneCycleLR	0.603	No gain — MultiScale not helpful here
v5 retrain (final)	Unique filename; 3 runs, best selected	0.577	Production model

6.3 Energy Estimation Experiments

Approach	MAE (ON)	kWh Error	Status
Direct regression (all windows)	280.4 W	100%	FAILED — predicts ~0W always
Mean watts (434.3 W)	413.3 W	37.6%	Used in production
ON-only regression	220.0 W	N/A ($R^2 < 0$)	FAILED — variance exceeds explainable signal
Seq2Seq (599-length output)	221.5 W	$\infty\%$	FAILED — model outputs explode to millions of watts

6.4 Final Model — washing_machine_v5_final.pt

Parameter	Value
Architecture	5-layer CNN-BiLSTM

Parameter	Value
Window size	599 minutes
Training strategy	Smart sampling: Houses 2 & 3 capped at 5K ON; Houses 5, 9, 11 full
F1 / Precision / Recall	0.577 / 0.561 / 0.593
Threshold	0.80
Power estimate	434.3 W mean ON watts

CHAPTER SEVEN: DISHWASHER — EXPERIMENTS & RESULTS

7.1 Challenge

Only House 5 provides meaningful dishwasher usage data (1.98% ON). The remaining four training houses each have ON ratios below 0.5%, making House 5 responsible for ~90% of all training signal. The model must generalise from one household's dishwasher pattern to House 20.

7.2 All Experiments

Experiment	F1	Threshold	Status
Exp 1: CNN-BiLSTM + BCE	0.560	0.85	Baseline
Exp 2: CNN-BiLSTM + Focal + Cosine LR	0.574	0.85	Slight improvement
Exp 3: Transformer Seq2Point	0.466	0.80	FAILED — flat probability distribution
Exp 4: CNN-BiLSTM + Optuna (40 trials)	0.661	0.70	Best single-scale model
Exp 5: CtRNN + BiGRU (Pirnat et al. 2024)	0.178	—	FAILED — wrong evaluation protocol in original paper
Exp 6: Data Augmentation ($\times 5$ ON)	0.385	—	FAILED — augmented distribution mismatch
Exp 7a: ResNet18 + Spectrogram	0.604	0.55	Best 2D approach
Exp 7b: ResNet18 + GAF	0.419	0.90	Weak
Exp 7c: ResNet18 + Recurrence Plot	0.329	0.65	Weak
Exp 8: MultiScale CNN-BiLSTM ★	0.726	0.80	WINNER
Exp 9: MultiScale + Optuna	0.714	0.70	Below Exp 8 — short-trial params do not generalise

7.3 The MultiScale Breakthrough

The winning approach changed the input, not the architecture. Each 599-point aggregate window becomes 4 channels:

Channel	Content	Purpose
ch0	Raw normalised aggregate	Direct power reading
ch1	First-order differential (normalised)	Detects cycle start/end

Channel	Content	Purpose
ch2	5-point moving average	Short-term local trend
ch3	15-point moving average	Long-term cycle context

This single change: +0.065 F1 (0.661 $\rightarrow$ 0.726), larger than any architecture change across all 9 experiments.

7.4 Final Model — dishwasher_multiscale_best.pt

Parameter	Value
Architecture	4-channel MultiScale CNN-BiLSTM
Window size	599 minutes
CNN layers	3 (channels: 4 $\rightarrow$ 16 $\rightarrow$ 128 $\rightarrow$ 128)
BiLSTM	hidden=256, 1 layer, bidirectional
Parameters	1,024,241
F1 / Precision / Recall	0.726 / 0.706 / 0.747
Threshold	0.80

CHAPTER EIGHT: CROSS-CUTTING LESSONS LEARNED

8.1 Window Size is the Most Impactful Hyperparameter

Choose a window containing 5–10 complete ON/OFF cycles of the target appliance. For the fridge (cycle ~62 min), 99 min is optimal. For the kettle (cycle ~1 min), 99 min provides sufficient context. For the washing machine (cycle 60–90 min), 599 min captures the full multi-phase cycle. Too-large a window buries local signal in irrelevant context.

8.2 Validation Strategy

Approach	Problem	Recommendation
Balanced 50/50 val sample	Inflated F1; wrong model selected as best	Never use for final evaluation
1:6 val sample (during training)	Still inflated; useful only for epoch-level feedback	OK during training epochs only
Full 623K val + threshold sweep	Honest; reflects real-world distribution	Always use for final reported results

8.3 Critical Bugs Fixed

Bug	Symptom	Fix
Double Sigmoid	Loss stuck at 0.52; probs collapse to 0.5–0.73	Remove nn.Sigmoid() from model; BCEWithLogitsLoss applies sigmoid internally
pos_weight on balanced data	Model predicts ON for everything from epoch 1	Remove pos_weight when training is 50/50; use Focal Loss
PyTorch 2.6 GradScaler	FutureWarning deprecation	Use torch.amp.GradScaler('cuda')
PyTorch 2.6 autocast	FutureWarning deprecation	Use torch.amp.autocast('cuda')
PyTorch 2.6 torch.load	UnpicklingError	Add weights_only=False
pandas resample '1T'	DeprecationWarning	Use '1min'
pandas fillna method	FutureWarning	Use series.ffill(limit=N)
Seq2Seq CNN size mismatch	HuberLoss crashes: 596 vs 599	Crop CNN output: f[:, :, :x.shape[2]]
Model overwritten by experiments	Original v5 model lost	Use unique filenames per experiment

8.4 Architecture Lessons

CNN alone outperforms CNN-BiLSTM for short-duration events (kettle). BiLSTM better for sustained appliances (fridge, washing machine, dishwasher). One BiLSTM layer often outperforms two layers (found by Optuna for dishwasher). MultiScale 4-channel input gave the largest single improvement in dishwasher (+0.065 F1) — more than any architecture change. Focal Loss ($\gamma \approx 2-3$, $\alpha \approx 0.6-0.75$) consistently outperforms BCE for imbalanced NILM tasks. Threshold sweep (0.10→0.99, step 0.05) is mandatory — optimal threshold varies from 0.22 (fridge) to 0.85 (kettle) and cannot be fixed in advance.

CHAPTER NINE: STREAMLIT APPLICATION — NILM_V2

9.1 Overview

The final system is deployed as a five-page Streamlit web application (nilm_v2) that loads all four production models and presents disaggregation results interactively. The app was validated against real House 20 ground truth data across a 3-day test window (4,320 minutes = 72 hours).

Page	Description
Disaggregation	Upload aggregate CSV → run all four models → display per-appliance ON/OFF timelines and energy summary
Analytics Dashboard	Historical energy intelligence: total kWh, appliance ranking, daily trends
Bill Estimator	Per-appliance kWh → cost using Jordan EMRC tiered residential tariff (April 2022 rates)
Forecasting	7-day energy forecast per appliance based on observed usage rate
Savings Simulator	Interactive sliders to simulate usage changes and calculate real-time bill impact

9.2 App Screenshots

Page 1 — Disaggregation: Aggregate Signal Upload






Figure 6.1: Aggregate signal upload — 4,320 readings (72h), mean load 460W, with Run Disaggregation button**Page 1 — Disaggregation: Results & Energy Breakdown****Figure 6.2:** Disaggregation results — 9.97 kWh total identified across 4 appliances with donut chart and bar chart breakdown**Page 1 — Disaggregation: Appliance Activity Timeline & Summary Table****Figure 6.3:** Per-appliance activity timeline (Kettle, Fridge, Washing Machine, Dishwasher) with appliance summary table showing ON time, energy, estimated bill, and architecture

Page 2 — Analytics Dashboard



Figure 6.4: Analytics Dashboard — overall statistics, appliance energy ranking (Fridge 5.74 kWh, WM 1.85 kWh, Dishwasher 1.51 kWh, Kettle 0.87 kWh) and daily energy trends

Page 3 — Bill Estimator




Figure 6.5: Electricity Bill Estimator — Jordan EMRC residential tariff, tariff tier breakdown, per-appliance monthly estimate in JOD

Page 4 — Energy Forecasting




Figure 6.6: Energy Forecasting — 7-day usage forecast per appliance with projected monthly bill of 3.891 JOD

Page 4 — Savings Simulator




Figure 6.7: Savings Simulator — interactive usage sliders per appliance with real-time bill impact calculation

9.3 Technical Implementation

All models loaded at startup with `torch.load(..., weights_only=False)` for PyTorch 2.6 compatibility. Per-chunk min-max normalisation on input signal — no global scaler, same as training. MultiScale input (for dishwasher) computed on-the-fly: raw + differential + $MA \times 5$ + $MA \times 15$. Threshold for each appliance loaded directly from model checkpoint.

Energy estimation: $\text{on_minutes} \times \text{mean_watts} / 60 / 1000$ for kWh; extrapolated to 30-day monthly estimate. Jordan EMRC tiered tariff: 33 fils/kWh (0–160), 76 fils/kWh (161–300), 114 fils/kWh (301–500), 141 fils/kWh (500+). Savings simulator: interactive sliders allow adjusting per-appliance usage minutes with real-time bill impact calculation. 7-day forecast: projects daily kWh per appliance based on observed usage rate from uploaded data.

CHAPTER TEN: REAL-WORLD TESTING & FINAL RESULTS

10.1 Test Setup

Parameter	Value
Test Period	3 days: 2014-09-29, 2014-11-17, 2014-10-20
Total Minutes	4,320 minutes (72 hours)
Validation House	House 20 (unseen during all training)
Input to System	Aggregate power only (Watts, 1-minute resolution)
Ground Truth	Individual appliance meter readings from House 20

10.2 Overall Results — 3 Days Combined

Appliance	Ground Truth (min)	Predicted (min)	Difference	Notes
Kettle	16	21	+5 min	Close estimate for a very short-duration appliance
Fridge	1,864	4,049	+2,185 min	Over-predicted — high recall, low threshold (0.22)
Washing Machine	223	139	−84 min	Precision 95% on Day 1; under-predicted on Days 2–3
Dishwasher	187	101	−86 min	Consistent with training F1=0.726 at threshold=0.80

10.3 Per-Day Breakdown

Date	Kettle	WM (Pred/GT)	Dishwasher (Pred/GT)	Highlight
2014-09-29	4 / 6 min	64 / 68 min	40 / 62 min	WM F1=0.91 — best single-day result
2014-11-17	7 / 6 min	0 / 68 min	35 / 0 min	Kettle F1=0.77 — accurate detection
2014-10-20	8 / 4 min	0 / 86 min	39 / 0 min	Fridge F1=0.66

10.4 Key Observations

Washing Machine — Best Performance: On 2014-09-29, the system predicted 64 ON minutes against a ground truth of 68 — a difference of only 4 minutes, corresponding to Precision of 95%. This is the strongest single-day result across all appliances.

Kettle — Accurate Detection: The kettle operates for only 2–6 minutes per use. On 2014-11-17, the system predicted 7 ON minutes against a ground truth of 6 — a near-perfect result. Across all 3 days, kettle total (21 min) is within 5 minutes of ground truth (16 min).

Fridge — High Recall, Over-Prediction: The fridge model consistently detected when the fridge was running (high recall) but over-predicted total ON time by approximately 2×. This is a direct consequence of the deliberately low classification threshold (0.22), set to maximise recall at the cost of precision — consistent with the F1=0.69 ceiling identified during training.

Dishwasher — Under-Prediction: The dishwasher was under-predicted across all three days. Dishwasher cycles are long but infrequent, and the model trained primarily on House 5 patterns.

10.5 Final Results vs Literature

Paper / System	Eval. Protocol	Kettle F1	Fridge F1	WM F1	Dishwasher F1
Liao et al. 2014 (IEEE)	Mixed splits, various datasets	~0.65	~0.55	~0.60	~0.58
Kelly & Knottenbelt 2015	Partially seen houses	~0.74	~0.60	~0.65	~0.62
Our Project	Fully unseen House 20 — STRICTEST	0.784	0.688	0.577	0.726

Our results use a strictly harder evaluation protocol than both reference papers. Despite this, kettle and dishwasher results exceed published values. The washing machine is slightly below published ranges, attributable to data quality issues in training houses.

10.6 Production Model Inventory

File	Appliance	F1	Window	Architecture
kettle_v8_w99.pt	Kettle	0.784	99 min	Pure CNN
fridge_v4_best.pt	Fridge	0.688	99 min	CNN-BiLSTM

File	Appliance	F1	Window	Architecture
washing_machine_v5_final.pt	Washing Machine	0.577	599 min	CNN-BiLSTM
dishwasher_multiscale_best.pt	Dishwasher	0.726	599 min	4-ch MultiScale CNN-BiLSTM

CHAPTER ELEVEN: CONCLUSIONS

This project delivered a complete end-to-end NILM system for household appliance energy disaggregation, from raw 8-second REFIT data to a deployed Streamlit web application. Four dedicated deep learning models were trained and deployed, providing disaggregation, analytics, cost estimation, forecasting, and a savings simulator.

The most significant technical contributions are:

Per-appliance window size optimisation: demonstrating that window size is more impactful than model architecture across all appliances tested.

Honest evaluation protocol: all results reported on completely unseen House 20 under strict LOHO — more rigorous than most published NILM work on the REFIT dataset.

MultiScale input representation: 4-channel input (raw + differential + two moving averages) improved dishwasher F1 by +0.065 — the largest single gain across any experiment.

Diagnostic tooling: ON/OFF probability separation as a leading indicator of model quality; threshold sweeping as a required step; per-appliance validation strategy calibration.

Complete failure documentation: double sigmoid bugs, pos_weight misuse, balanced validation inflation, Transformer failures, data augmentation limitations, and regression head conflicts — providing a rigorous basis for future NILM research.

The Streamlit application demonstrates that the trained models can power a practical energy monitoring tool, providing actionable insights on appliance usage, cost

breakdown, and savings simulation — directly contributing to energy efficiency goals for residential users.

Future Work

Training at native 8-second resolution (rather than 1-minute resampled) improved kettle F1 to 0.91 in a supplementary experiment; extending this to all appliances is the highest-priority next step. Cross-dataset transfer learning on a Jordanian household dataset would improve deployment relevance. The Microwave (F1=0.40 ceiling) remains unsolved; event-triggered detection may prove more effective. Integration with a live smart meter API would enable real-time disaggregation rather than offline file upload.

REFERENCES

- [1] Hart, G. W. (1992). Nonintrusive appliance load monitoring. *Proceedings of the IEEE*, 80(12), 1870–1891.
- [2] Kelly, J. and Knottenbelt, W. (2015). Neural NILM: Deep neural networks applied to energy disaggregation. *Proceedings of the 2nd ACM BuildSys Conference*, Seoul, South Korea.
- [3] Liao, J., Elafoudi, G., Stankovic, L. and Stankovic, V. (2014). Non-intrusive appliance load monitoring using low-resolution data. *IEEE SmartGridComm*, Venice, Italy.
- [4] Murray, D., Stankovic, L. and Stankovic, V. (2017). An electrical load measurements dataset of United Kingdom households from a two-year monitoring study. *Scientific Data*, 4, 160122.
- [5] Lin, T. Y., Goyal, P., Girshick, R., He, K. and Dollar, P. (2017). Focal loss for dense object detection. *IEEE ICCV*, Venice, Italy.
- [6] Bonfigli, R. et al. (2018). Non-intrusive load monitoring by using active and reactive power in additive Factorial Hidden Markov Models. *Applied Energy*, 208, 1590–1607.
- [7] Nalmpantis, C. and Vrakas, D. (2019). Machine learning approaches for non-intrusive load monitoring. *Artificial Intelligence Review*, 52(1), 217–243.
- [8] Pirnat, A., Lustrek, M. and Gradisek, A. (2024). NILM with continuous-time recurrent neural networks. *Energy and Buildings*, 302, 113712.
- [9] Electricity and Minerals Regulatory Commission (EMRC), Jordan. (2022). *Electricity Tariffs for Residential Customers*. Amman: EMRC.